

Heart Attack Risk Prediction: A Case Study

This Data is regarding the Heart Attack Risk Prediction where we having the multiple features to check the Heart Attack Risk. Now we'll perform a lot of analysis to our dataset also we'll create a ML model which will help us to get to know which Data consist to more heart attack risk.



About the Dataset:

- **Age-:** Age of the Patient
- **Sex-:** Gender of the Patient
- **CP-:** Type of Chest Pain
 - Value 1: typical angina
 - Value 2: A-typical angina
 - Value 3: non-anginal pain
 - Value 4: asymptomatic
- **trtbps-:** trestbps (Resting Blood Pressure)
- **Chol-:** Cholesterol
- **FBS-:** Fasting Blood Sugar
- **rest_ecg-:** Resting Electrocardiographic Results
 - Value 0: normal
 - Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV)
 - Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria
- **exang-:** Exercise Induced Angina (1 - Yes; 0 - No)
- **OldPeak-:** Previous Peak
- **SLP-:** Slope
- **caa-:** Number of Major Vessels
- **thall-:** Thallium Stress Test Result (0-3)
- **Output-:** Whether the Patient Has Heart Attack Risk or Not
 - 1 - The Person Has Heart Attack Risk
 - 2 - The Person Does Not Have Heart Attack Risk

iatric

Loading the Standard Libraries

In [190...

```
import numpy as np
import pandas as pd
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as sts
```

Loading the Dataset

```
In [191...] df = pd.read_csv('Dataset/heart.csv')
df.head()
```

```
Out[191...]   age  sex  cp  trtbps  chol  fbs  restecg  thalachh  exng  oldpeak  slp  caa  thall  output
0    63    1   3   145   233    1     0     150     0     2.3    0   0    1     1
1    37    1   2   130   250    0     1     187     0     3.5    0   0    2     1
2    41    0   1   130   204    0     0     172     0     1.4    2   0    2     1
3    56    1   1   120   236    0     1     178     0     0.8    2   0    2     1
4    57    0   0   120   354    0     1     163     1     0.6    2   0    2     1
```

Fetching some information with data

```
In [192...] df.shape
```

```
Out[192...] (303, 14)
```

```
In [193...] df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null    int64
 1   sex         303 non-null    int64
 2   cp          303 non-null    int64
 3   trtbps      303 non-null    int64
 4   chol        303 non-null    int64
 5   fbs         303 non-null    int64
 6   restecg     303 non-null    int64
 7   thalachh    303 non-null    int64
 8   exng        303 non-null    int64
 9   oldpeak     303 non-null    float64
10   slp         303 non-null    int64
11   caa         303 non-null    int64
12   thall       303 non-null    int64
13   output      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

Observation-: There having only **303** Record with **14** Features also there is no as such object or any irrelevant data type format which we need to transform all data are in the form of **int** and **float**

In [194... `df.describe().T`

Out[194...

	count	mean	std	min	25%	50%	75%	max
age	303.0	54.366337	9.082101	29.0	47.5	55.0	61.0	77.0
sex	303.0	0.683168	0.466011	0.0	0.0	1.0	1.0	1.0
cp	303.0	0.966997	1.032052	0.0	0.0	1.0	2.0	3.0
trtbps	303.0	131.623762	17.538143	94.0	120.0	130.0	140.0	200.0
chol	303.0	246.264026	51.830751	126.0	211.0	240.0	274.5	564.0
fbs	303.0	0.148515	0.356198	0.0	0.0	0.0	0.0	1.0
restecg	303.0	0.528053	0.525860	0.0	0.0	1.0	1.0	2.0
thalachh	303.0	149.646865	22.905161	71.0	133.5	153.0	166.0	202.0
exng	303.0	0.326733	0.469794	0.0	0.0	0.0	1.0	1.0
oldpeak	303.0	1.039604	1.161075	0.0	0.0	0.8	1.6	6.2
slp	303.0	1.399340	0.616226	0.0	1.0	1.0	2.0	2.0
caa	303.0	0.729373	1.022606	0.0	0.0	0.0	1.0	4.0
thall	303.0	2.313531	0.612277	0.0	2.0	2.0	3.0	3.0
output	303.0	0.544554	0.498835	0.0	0.0	1.0	1.0	1.0

Observation:-

- Age :-
 - The Maximum Age of Our Patient is :- 77
 - The Miniumu Age of Our Patient is :- 29
 - The Average Age of Our Patient is :- 55
- trtbps : - **Resting Blood Presure**

Our Resing Blood presure reading should be between 120 / 180. if it's more than 180 it's very critical condition we should immediate seek for medical treatment * The Maximum trtbps of Our Patient is :- 200 which is very critical condition for our patients * The Minimum trtbps of Our

Patient is :- 94 which is low blood pressure which is not critical but it's not that much good for our health * The Average trtbps of Our Patients is :- 131 which is pretty normal

- Chol :- **Cholesterol**

Our Cholesterol Level should be under 200-240 * The Maximum Cholesterol of Our Patient is :- 564 which can significantly increase the risk of heart attack * The minimum cholesterol of Our Patient is :- 211 which can also significantly increase the risk of heart attack * The Average Cholesterol of Our Patient is :- 246 which is pretty normal

Let's do some Sanity of the Data

```
In [195... df.duplicated().sum()
```

```
Out[195... 1
```

```
In [196... df[df.duplicated()]
```

```
Out[196... 
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
164	38	1	2	138	175	0	1	173	0	0.0	2	4	2	1

Observation:- Our Data is containing only 1 duplicate values let's drop it out

```
In [197... df = df.drop_duplicates()
```

Let's Check the Unique Value in our dataset

```
In [198... df.nunique()
```

```
Out[198... age      41
sex        2
cp         4
trtbps     49
chol      152
fbs        2
restecg    3
thalachh   91
exng       2
oldpeak    40
slp        3
caa        5
thall      4
output     2
dtype: int64
```

Let's Check the Null Value From Our Dataset

```
In [199... df.isnull().sum()
```

```
Out[199... age      0
sex      0
cp       0
trtbps   0
chol     0
fbs      0
restecg  0
thalachh 0
exng     0
oldpeak  0
slp      0
caa      0
thall    0
output   0
dtype: int64
```

Observation:- There are no as such missing values in our dataset


EDA

```
In [200... df['sex'].value_counts()
```

```
Out[200... sex
1    206
0     96
Name: count, dtype: int64
```


```
In [201... cate_value = ['sex', 'cp', 'fbs', 'restecg', 'exng', 'slp', 'caa', 'thall']
```

```
In [202... fig , axes = plt.subplots(2,4,figsize = (18,11))
axes = axes.flatten()
for i in range(len(cate_value)):
    sns.countplot(data = df , x = cate_value[i], ax = axes[i])
```

Countplot of Target

```
In [203... sns.countplot(data = df , x = 'output')
plt.show()
```



Let's See the Correlation on the matrix

```
In [204... num_col = ['age', 'trtbps', 'chol', 'thalachh', 'oldpeak']  
df_corr = df[num_col].corr().T  
df_corr
```







Out[204...

	age	trtbps	chol	thalachh	oldpeak
age	1.000000	0.283121	0.207216	-0.395235	0.206040
trtbps	0.283121	1.000000	0.125256	-0.048023	0.194600
chol	0.207216	0.125256	1.000000	-0.005308	0.050086
thalachh	-0.395235	-0.048023	-0.005308	1.000000	-0.342201
oldpeak	0.206040	0.194600	0.050086	-0.342201	1.000000

Outliers

In [205...

```
fig , axes = plt.subplots(3,2 , figsize = (18,14))
axes = axes.flatten()
for i in range(len(num_col)):
    sns.boxplot(data = df , x = num_col[i], ax = axes[i])
plt.show()
```





Obsevation--: As we can see there are some outliers in our dataset we can remove this if we don't lose a lot of data

Removing the Outliers






```
In [206... s = set()
for i in df[num_col]:
    q1 = df[i].quantile(0.25)
    q3 = df[i].quantile(0.75)
    iqr = q3-q1
    lower_bond = q1 - 1.5 * (iqr)
    upper_bond = q3 + 1.5 * (iqr)
    index_ = df[(df[i] < lower_bond ) | (df[i] > upper_bond)].index.tolist()
    s.update(index_)
print(f"The Index Where We Have the Outliers : {s}")
print(f"The Number of Data Which we lost : {len(s)}")
df = df.drop(s)
```


The Index Where We Have the Outliers : {260, 8, 266, 272, 28, 291, 203, 204, 85, 220, 221, 223, 96, 101, 110, 241, 246, 248, 250}

The Number of Data Which we lost : 19

Observation:- We have lost almost 19 records which were treated as outliers now if we check that our dataset is still containing the outliers after applying the method of removing the outliers

```
In [207... fig , axes = plt.subplots(3,2 , figsize = (18,14))
axes = axes.flatten()
for i in range(len(num_col)):
    sns.boxplot(data = df , x = num_col[i], ax = axes[i])
plt.show()
```











Observation:- As you can see there are no as such outliers in our dataset

Now We'll See the Distribution of the Continuous Column

```
In [208... fig, axes = plt.subplots(3,2,figsize = (18,14))
axes = axes.flatten()
for i in range(len(num_col)):
    sns.histplot(data = df, x = num_col[i], hue = 'output', ax = axes[i], fill = True, kde = True)
```



oldpeak

Observation-: As we can see our age , trtbp , chol are normally distributed whereas oldpeak is left-skewed distribution and thalachh is a bit right skewed distributed.

- **Age:-**
 - More Risk of Heart Attack is between to Age Group of 40-55
 - Whereas is the Less Risk of Heart Attack is Between Age Group of 55-65
- **trtbps:- trestbps (Resting Blood Pressure).**
 - More Risk of Heart Attack Patient with low BP 95- 105 also with high BP of 120-140 some there is some patient is also there with the BP of 145-155.
 - Where as Less Risk of Heart Attack Patient **Resting Blood Pressure** is 140-145 also there some patient has been observed with low Risk of Heart Attack with the BP of 155-170
- **Chol (Cholesterol) :-.**
 - Cholesterol Level Make Huges impact of Heart Attack Risk. More People Having Heart Attack Risk with the Cholesterol Level of 145-250. but few people also have been observed with Heart Attack Risk at Cholesterol Level 350.
 - Where Less Risk of Heart Attack Risk of Patient has been observed 260 - 300.
- **thalachh (The Person Heart Rate Achieved) :-**
 - The Higher Risk of Heart Attack Patinent is 150-200
 - The Low Risk of Heart Attack Patient is 90 - 140

Data Preprocessing

1) Seperating the Dependent and Independent Features

```
In [209... df_ = df.copy()
x = df_.drop('output', axis = 1)
y = df_[['output']]
```

2) Splitting the Dataset

```
In [210... from sklearn.model_selection import train_test_split
```

```
In [211... x_train , x_test , y_train, y_test = train_test_split(x,y,random_state = 42 , test_size = 0.25)
```

```
In [212... x_train.shape , y_train.shape , x_test.shape , y_test.shape
```

```
Out[212... ((212, 13), (212, 1), (71, 13), (71, 1))
```

Feature Engineering

1) Numerical Data Encoding

```
In [213... from sklearn.preprocessing import StandardScaler
```

```
In [214... ss = StandardScaler()  
ss
```

```
Out[214... ▾ StandardScaler  
StandardScaler()
```

```
In [215... x_train[['age', 'trtbps', 'chol', 'thalachh']] = ss.fit_transform(x_train[['age', 'trtbps', 'chol', 'thalachh']])
```

```
In [216... def transform(x_test):  
    x_test[['age', 'trtbps', 'chol', 'thalachh']] = ss_new.transform(x_test[['age', 'trtbps', 'chol', 'thalachh']])  
    return x_test
```

```
In [217... import pickle  
f = open("Feature Scaler File\\StandardScaler.pkl", "wb")  
pickle.dump(ss, f)  
f.close()
```

Modeling

1) LogisticRegression

```
In [218... from sklearn.linear_model import LogisticRegression
```

```
In [219... lr = LogisticRegression()  
lr
```

```
Out[219... ▼ LogisticRegression  
LogisticRegression()
```

```
In [220... lr.fit(x_train , y_train)
```

```
Out[220... ▼ LogisticRegression  
LogisticRegression()
```

```
In [221... lr.coef_
```

```
Out[221... array([[ 0.02512152, -1.30496145,  0.67026635, -0.28378598, -0.35751672,  
         -0.28348342,  0.63866043,  0.74933245, -0.74104964, -0.54843962,  
         0.49907719, -0.5671385 , -1.08789224]])
```

```
In [222... lr.intercept_
```

```
Out[222... array([3.1373863])
```

```
In [223... import pickle
```

```
In [224... file = open("Feature Scaler File\\StandardScaler.pkl", "rb")  
ss_new = pickle.load(file)
```

```
In [225... def transform(x_test):  
    x_test[['age', 'trtbps', 'chol', 'thalachh']] = ss_new.transform(x_test[['age', 'trtbps', 'chol', 'thalachh']])  
    return x_test  
  
x_test_ = transform(x_test)
```

```
In [226... y_pred_test = lr.predict(x_test_)
```

```
In [227... pd.DataFrame({"predicted":y_pred_test, "actual":y_test['output']}).head()
```

```
Out[227... 
```

	predicted	actual
10	1	1
263	1	0
145	1	1
216	0	0
77	1	1

```
In [228... lr.score(x_train, y_train)
```

```
Out[228... 0.8632075471698113
```

Accuracy

```
In [229... from sklearn.metrics import classification_report , accuracy_score , confusion_matrix
```

```
In [230... print(classification_report(y_test ,y_pred_test ))
print(confusion_matrix(y_test, y_pred_test))
print(f"Accuracy Score of the Test Data-: {round(accuracy_score(y_test , y_pred_test)*100,3)} %")
```

	precision	recall	f1-score	support
0	0.81	0.70	0.75	30
1	0.80	0.88	0.84	41
accuracy			0.80	71
macro avg	0.80	0.79	0.79	71
weighted avg	0.80	0.80	0.80	71

```
[[21  9]
 [ 5 36]]
```

Accuracy Score of the Test Data-: 80.282 %

Observation:- We got Test Accuracy as 83.099 we'll perform the hyperparameter tuning to see weather we can improve the accuracy of the model or not

Hyperparameter Tunning

```
In [231... from sklearn.model_selection import GridSearchCV
```

```
In [232... parameter = {"penalty":["l1","l2","elasticnet"], "C" : [10,50,100, 500], "random_state" : [42], "solver":["liblinear","saga"]}
```

```
In [233... gsv = GridSearchCV(lr , param_grid = parameter, cv = 10)
```

```
In [234... gsv
```

```
Out[234... 
└─ GridSearchCV
  └─ estimator: LogisticRegression
    └─ LogisticRegression
```

```
In [235... np.ravel(y_train).shape
```

```
Out[235... (212,)
```

```
In [236... import warnings
warnings.filterwarnings('ignore')
gsv.fit(x_train ,y_train['output'])
```

```
Out[236... GridSearchCV
  estimator: LogisticRegression
    LogisticRegression
```

```
In [237... gsv.best_params_
```

```
Out[237... {'C': 10, 'penalty': 'l2', 'random_state': 42, 'solver': 'liblinear'}
```

```
In [238... gsv.best_estimator_
```

```
Out[238... LogisticRegression
LogisticRegression(C=10, random_state=42, solver='liblinear')
```

```
In [239... lr_ = LogisticRegression(C= 10, penalty='l2', random_state= 42, solver='liblinear')
```

Observation-: We have Added the Best Predicted Parameter into Logistic Regression. Now We'll See is there any changes in Accuracy Score

```
In [240... lr_.fit(x_train, y_train)
```

```
Out[240... LogisticRegression
LogisticRegression(C=10, random_state=42, solver='liblinear')
```

```
In [241... lr_.coef_
```

```
Out[241... array([[ 0.05318007, -1.60969064,  0.72833798, -0.32335647, -0.40054777,
        -0.36606581,  0.7601775 ,  0.82466617, -0.84674224, -0.52783103,
         0.6169009 , -0.58740307, -1.16395227]])
```

In [242... lr_.intercept_

Out[242... array([3.27403243])

In [243... y_pred_test_ = lr_.predict(x_test)

In [244... print(f"The Training Score : {lr_.score(x_train , y_train)*100}")

The Training Score : 86.79245283018868

Now we'll check the score of accuracy

In [245... print(classification_report(y_test ,y_pred_test_))
 print(confusion_matrix(y_test, y_pred_test_))
 print(f"Accuracy Score of the Test Data-: {round(accuracy_score(y_test , y_pred_test_)*100,3)} %")

	precision	recall	f1-score	support
0	0.75	0.70	0.72	30
1	0.79	0.83	0.81	41
accuracy			0.77	71
macro avg	0.77	0.76	0.77	71
weighted avg	0.77	0.77	0.77	71

```
[[21  9]
 [ 7 34]]
```

Accuracy Score of the Test Data-: 77.465 %

As we can see we have improvement of 1% of the Model

In [246... f = open("ML Model//Logistic_Regression.pkl", "wb")
 pickle.dump(lr_, f)
 f.close()

DecisionTree Classifier

Now We'll Try Another Algorithm Which is know as Decision

```
In [247... from sklearn.tree import DecisionTreeClassifier
```

```
In [248... dtc = DecisionTreeClassifier()  
dtc
```

```
Out[248... ▼ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [249... dtc.fit(x_train , y_train)
```

```
Out[249... ▼ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [250... dtc_y_pred = dtc.predict(x_test)
```

```
In [251... print(classification_report(y_test ,dtc_y_pred ))  
print(confusion_matrix(y_test, dtc_y_pred))  
print(f"Accuracy Score of the Test Data-: {round(accuracy_score(y_test ,dtc_y_pred)*100,3)} %")
```


	precision	recall	f1-score	support
0	0.60	0.70	0.65	30
1	0.75	0.66	0.70	41
accuracy			0.68	71
macro avg	0.68	0.68	0.67	71
weighted avg	0.69	0.68	0.68	71

```
[[21  9]
 [14 27]]
```

Accuracy Score of the Test Data-: 67.606 %

```
In [252... dtc.get_depth()
```

```
Out[252... 9
```

```
In [253... dtc = DecisionTreeClassifier(max_depth = 3)
dtc.fit(x_train , y_train)
```

```
Out[253... ▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3)
```

```
In [254... print(f"The Accuracy of Train Dataset is -: {dtc.score(x_train , y_train) * 100}")
```

The Accuracy of Train Dataset is -: 85.84905660377359

```
In [255... dtc_y_pred_test = dtc.predict(x_test)
dtc_y_pred_test
```

```
Out[255... array([1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1,
        1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1,
        1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1,
        1, 1, 0, 1, 1], dtype=int64)
```

```
In [256... print(classification_report(y_test ,dtc_y_pred_test ))
print(confusion_matrix(y_test, dtc_y_pred_test))
print(f"Accuracy Score of the Test Data-: {round(accuracy_score(y_test ,dtc_y_pred_test)*100,3)} %")
```

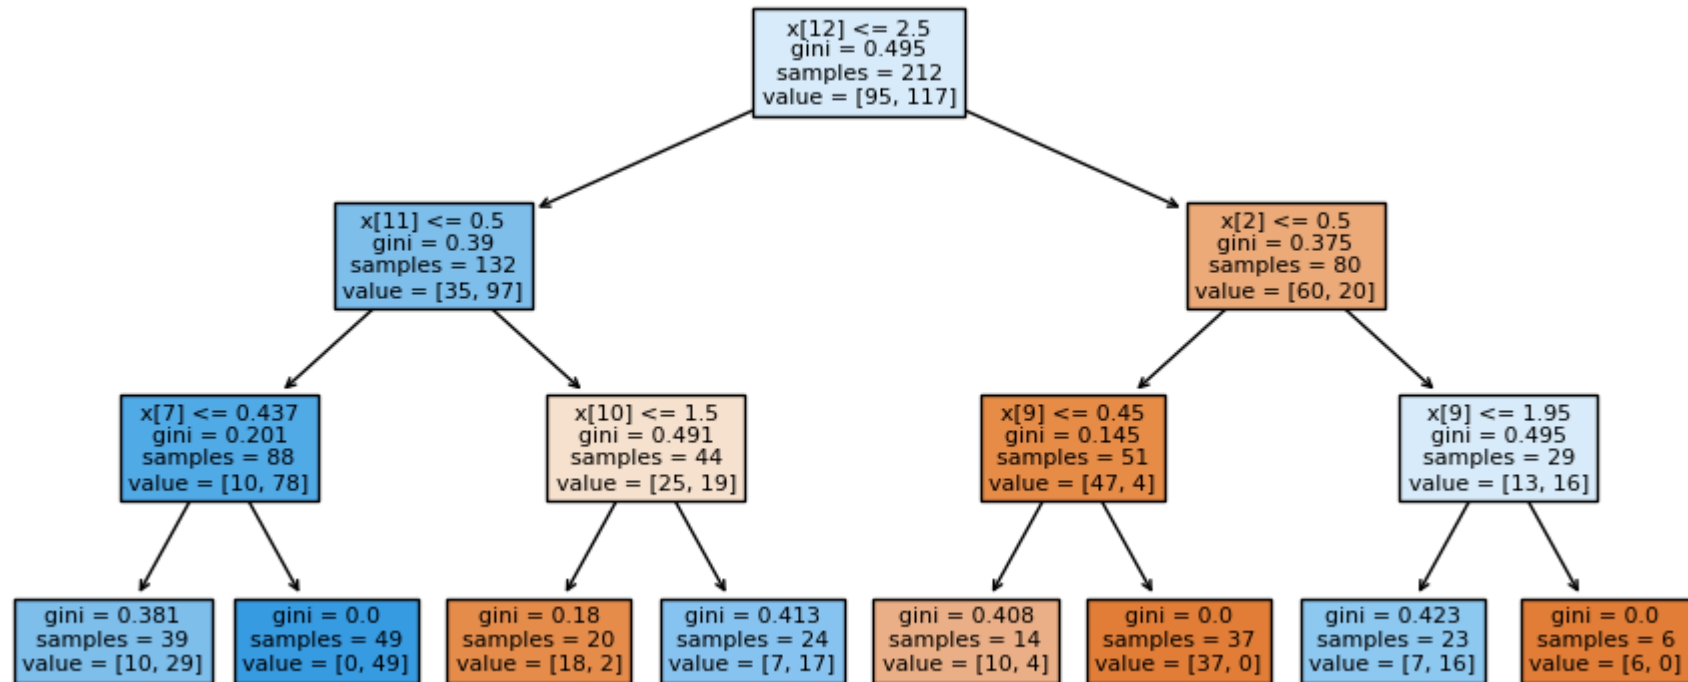
	precision	recall	f1-score	support
0	0.79	0.73	0.76	30
1	0.81	0.85	0.83	41
accuracy			0.80	71
macro avg	0.80	0.79	0.80	71
weighted avg	0.80	0.80	0.80	71

```
[[22  8]  
 [ 6 35]]
```

Accuracy Score of the Test Data-: 80.282 %

```
In [257... from sklearn.tree import plot_tree
```

```
In [258... plt.figure(figsize = (11,5))  
plot_tree(dtc, filled =True)  
plt.show()
```



```
In [264... f = open("ML Model/DecisionTree_Model.yml", "wb")
pickle.dump(dtc ,f )
f.close()
```

Accuracy

- **LogisticRegression-:** We get accuracy with using LogisticRegression of Training Dataset is :- **86.73%** and test Dataset with score of **77.46%**
- **DecisionTree Classifier-:** We Get Accuracy With Using DecisionTree Classifier of Training Dataset is :- **85.84%** and test Dataset with score of **80.28%**

So we use **DecisionTree Classifier** for model **DecisionTree Classifier**.